

Content extraction and search using Apache Tika

—

Employment Postings

Dataset

contributed via DARPA XDATA

Renu Kankamedala

Aldrin Rodrigues

Neha Ahuja

Poushali Banerjee

CSCI 572 HW 1

Prof. Chris Mattman

October 6, 2014

University of Southern California

Table of Contents:

- Introduction
- Methodology
- Results

other TODO's

1. Change hard coded file paths in DeduplicationJSON. (Aldrin)
2. Create a readme.txt with any instructions for grader, any external libraries (Renu)
3. Deduplication - used Jaro Winkler Similarity measure for Location2 column - java code for that function is taken from code.google.com

Introduction

In this project we create a local search engine of the Employment postings by cleansing, transforming, and developing an algorithm for ranking the job postings from <http://www.computrabajo.com> affiliate sites that primarily serve Mexico and South American countries. We have worked with a subset of the data for the most except where mentioned.

Methodology

The data is given in tsv format which is messy and difficult to understand, so first we transform the TSV files into Java Script Object Notation (JSON) files, and then from there into individual JSON posting files per job.

The TSV Parser:

In the parser, we reach each line and extract the column values. We then use the `WriteoutContentHandler(-1)` to get rid of 1000000 limit for writing characters and `xhtml` content handler for creating `xhtml` tags. To start a `xhtml` document we used `xhtml.startDocument()` which by default creates the namespace, `html` and `body` tags. Then we create a table using `xhtmlcontenthandler's startElement` and `endElement` functions to create different table tags. We used `.characters()` function to print the values and the `writer` function to write the `xhtml` output to a file.

TSV to Json:

Once we have the `xhtml` file we used `html cleaner` to parse through the `html` tags and retrieve the values. Similar to `XHTMLContentHandler` we created `JsonContentHandler` and overrided the functions like `startDocument`, `endDocument`, `startEle`, `endEle` accordingly and added few new methods like `startArray`, `startJsonAttribute` for different elements in `Json`. Thus we created `JSON` files corresponding to the `XHTML` table rows (one file per row).

Crawler:

After we had all the components for translating `tsv` files into `JSON`, we developed a program which fetches all the `TSV` files from the respective folder and converts to `xhtml` using `TSV` parser and the `xhtml` content is passed to `Json` content handler to create `json` files.

ETLLIB Crawler:

This is an API developed by Chris Mattmana and team to munge and prepare `JSON`, `TSV` and other data using `Apache Tika` and `JSON` parsing/loading for ETL via `Apache OODT` (or other libs) into `Apache Solr`. This was already provided to us and we installed the library without `tika` and ran the necessary commands for translating `tsv` to `json`.

We also wrote some code in `python` that would iterate through `JSON` files and count the number of files produced. Details about how we used the `etllib` and other code to achieve our results are explained in `readme.txt` (attached).

Deduplication:

With respect to web crawling, de-duplication refers to the identification of identical and nearly identical web pages and indexing only a single version to return as a search result. In this part we attempted to

remove any duplicate file/records from our JSON results. For this we followed the following methodology:

First, by analyzing the dataset we decide columns like Company, Department, Location2, Url, Title deduce uniqueness to a record. The same company can have a requirement in other department They can have requirement for same company and department in different locations and we used Jaro Winkler string similarity measure to compare two locations. Jaro winkler is good when the two fields have their prefix matching. Ex: San Francisco and San Francisco City. We used Jaro Winkler Similarity measure for Location2 column - java code for that function is taken from code.google.com Title, the Name of the position and Url - If the url's match we need not check for match between any other columns.

Algorithm:

```
if company name matches {
    if url matches {
        return true;
    } else if dept name matches {
        if title matches {
            if location2 columns are not null {
                if (JaroWinkler similarity score >= 0.8) {
                    return true;
                }
            } else if location2 column is null in both records {
                return true;
            }
        }
    }
}
```

Results

- Count of Json Jobfiles from our TSVParse: 12,483,611
- Count of Json Jobfiles from ETLLIB software: 12483611
- Deduplication (explained above)

- Number of job files produced with deduplication enabled:
- Number of job files produced without deduplication when run on the entire data set (extra credit): 119423497

Analysis
